

# Spotify

March 12, 2024

## 1 Machine Learning

### 1.0.1 Course-End Project Problem Statement

## Course-End Project: Creating Cohorts of Songs

### 1.0.2 Problem Scenario:

The customer always looks forward to specialized treatment, whether shopping over an e-commerce website or watching Netflix. They want what they might like to see. To keep the customers engaged, it is also crucial for companies to always present the most relevant information. Spotify is a Swedish audio streaming and media service provider. The company has over 456 million active monthly users, including over 195 million paying subscribers, as of September 2022. The company intends to create cohorts of different songs that will aid in the recommendation of songs to users based on various relevant features. Each cohort would contain similar types of songs. ### Problem Objective: As a data scientist, you should perform exploratory data analysis and perform cluster analysis to create cohorts of songs. The goal is to gain a better understanding of the various factors that contribute to creating a cohort of songs.

- Note: Download Data Dictionary – Creating cohorts of songs.xlsx from the course resource section in the LMS.

**Data Description:** This dataset contains data from Spotify’s API about all albums for the Rolling Stones listed on Spotify. It is important to note that all songs have unique IDs.

### Variable Description

- name It is the name of the song.
- album It is the name of the album.
- release\_date It is the day, month, and year the album was released.
- track number It is the order the song appears on the album.
- id It is the Spotify id for the song.
- uri It is the Spotify URI for the song.
- acousticness A confidence measure from 0.0 to 1.0 of whether the track is acoustic. 1.0 represents high confidence the track is acoustic.

- **danceability** It describes how suitable a track is for dancing based on a combination of musical elements, including tempo, rhythm stability, beat strength, and overall regularity. A value of 0.0 is the least danceable, and 1.0 is the most danceable.
- **energy** It is a measure from 0.0 to 1.0 and represents a perceptual measure of intensity and activity. Typically, energetic tracks feel fast, loud, and noisy. For example, death metal has high energy, while a Bach prelude scores low on the scale. Perceptual features contributing to this attribute include dynamic range, perceived loudness, timbre, onset rate, and general entropy.
- **instrumentalness** It predicts whether a track contains no vocals. “Ooh” and “aah” sounds are treated as instrumental in this context. Rap or spoken word tracks are clearly “vocal”. The closer the instrumentalness value is to 1.0, the greater likelihood the track contains no vocal content. Values above 0.5 are intended to represent instrumental tracks, but confidence is higher as the value approaches 1.0.
- **liveness** It detects the presence of an audience in the recording. Higher liveness values represent an increased probability that the track was performed live. A value above 0.8 provides a strong likelihood that the track is live.
- **loudness** The overall loudness of a track in decibels (dB) and loudness values are averaged across the entire track and are useful for comparing the relative loudness of tracks. Loudness is the quality of a sound that is the primary psychological correlate of physical strength (amplitude). Values typically range between -60 and 0 dB.
- **speechiness** It detects the presence of spoken words in a track. The more exclusively speech-like the recording (e.g., talk show, audiobook, poetry), the closer to 1.0 the attribute value. Values above 0.66 describe tracks that are probably made entirely of spoken words. Values between 0.33 and 0.66 describe tracks that may contain both music and speech, either in sections or layered, including such cases as rap music. Values below 0.33 most likely represent music and other non-speech-like tracks.
- **tempo** The overall estimated tempo of a track in beats per minute (BPM), and in musical terminology, the tempo is the speed or pace of a given piece and derives directly from the average beat duration.
- **valence** A measure from 0.0 to 1.0 describes the musical positiveness conveyed by a track, and tracks with high valence sound more positive (e.g., happy, cheerful, euphoric), while tracks with low valence sound more negative (e.g., sad, depressed, angry).
- **popularity** It is the popularity of the song from 0 to 100.
- **duration\_ms** It is the duration of the track in milliseconds.

## Steps to Perform:

**Initial data inspection and data cleaning:** Check whether the data has duplicates, missing values, irrelevant (erroneous entries) values, or outliers.

Depending on your findings, clean the data for further processing. ##### **Perform Exploratory Data Analysis and Feature Engineering:** Use appropriate visualizations to find out which two albums should be recommended to anyone based on the number of popular songs in an album.

Perform exploratory data analysis to dive deeper into different features of songs and identify the pattern.

Discover how a song's popularity relates to various factors and how this has changed over time. Comment on the importance of dimensionality reduction techniques, share your ideas and explain your observations. ##### Perform Cluster Analysis: Identify the right number of clusters Use appropriate clustering algorithm Define each cluster based on the features

```
[1]: # Imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.express as px
from plotly.subplots import make_subplots
import plotly.graph_objects as go
import warnings
```

```
[2]: # Suppressing the warnings
warnings.filterwarnings('ignore')
```

```
[3]: import scipy.cluster.hierarchy as sch
from sklearn.cluster import AgglomerativeClustering
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
```

## 1.1 Data Inspection and Data Cleaning

```
[4]: # Load Data
df = pd.read_excel('1673873388_rolling_stones_spotify.xlsx')
df.head()
```

```
[4]: Unnamed: 0      name      album release_date \
0      0  Concert Intro Music - Live  Licked Live In NYC  2022-06-10
1      1  Street Fighting Man - Live  Licked Live In NYC  2022-06-10
2      2      Start Me Up - Live  Licked Live In NYC  2022-06-10
3      3  If You Can't Rock Me - Live  Licked Live In NYC  2022-06-10
4      4      Don't Stop - Live  Licked Live In NYC  2022-06-10

      track_number      id      uri \
0      1  2IEkywLJ4ykbhi1yRQvmsT  spotify:track:2IEkywLJ4ykbhi1yRQvmsT
1      2  6GVgVJBKkGJoRfarYRvGTU  spotify:track:6GVgVJBKkGJoRfarYRvGTU
2      3  1Lu761pZ0dBTGpzzaQoZNW  spotify:track:1Lu761pZ0dBTGpzzaQoZNW
3      4  1agTQzOTUnGNgyckEqiDH  spotify:track:1agTQzOTUnGNgyckEqiDH
4      5  7piGJR8YndQBQWVXv6KtQw  spotify:track:7piGJR8YndQBQWVXv6KtQw

      acousticness  danceability  energy  instrumentalness  liveness  loudness \
0      0.0824      0.463      0.993      0.996000      0.932      -12.913
1      0.4370      0.326      0.965      0.233000      0.961      -4.803
```

2	0.4160	0.386	0.969	0.400000	0.956	-4.936
3	0.5670	0.369	0.985	0.000107	0.895	-5.535
4	0.4000	0.303	0.969	0.055900	0.966	-5.098

	speechiness	tempo	valence	popularity	duration_ms
0	0.1100	118.001	0.0302	33	48640
1	0.0759	131.455	0.3180	34	253173
2	0.1150	130.066	0.3130	34	263160
3	0.1930	132.994	0.1470	32	305880
4	0.0930	130.533	0.2060	32	305106

```
[5]: # Creating a backup dataframe and dropping unnecesssary solumms
df_bak = df.copy(deep=1)

df.drop(['Unnamed: 0', 'id', 'uri', 'duration_ms', 'track_number'], axis = 1,
        inplace = True)
```

```
[6]: print(f'>>> Shape : {df.shape}\n')
print(f'>>> Info : ')
df.info()
```

```
>>> Shape : (1610, 13)
```

```
>>> Info :
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1610 entries, 0 to 1609
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                  1610 non-null   object
1   album                 1610 non-null   object
2   release_date          1610 non-null   datetime64[ns]
3   acousticness          1610 non-null   float64
4   danceability           1610 non-null   float64
5   energy                 1610 non-null   float64
6   instrumentalness       1610 non-null   float64
7   liveness              1610 non-null   float64
8   loudness              1610 non-null   float64
9   speechiness           1610 non-null   float64
10  tempo                  1610 non-null   float64
11  valence                1610 non-null   float64
12  popularity             1610 non-null   int64
dtypes: datetime64[ns](1), float64(9), int64(1), object(2)
memory usage: 163.6+ KB
```

```
[7]: df.describe().transpose()
```

```
[7]:
```

	count	mean	std	min	25% \
acousticness	1610.0	0.250475	0.227397	0.000009	0.058350
danceability	1610.0	0.468860	0.141775	0.104000	0.362250
energy	1610.0	0.792352	0.179886	0.141000	0.674000
instrumentalness	1610.0	0.164170	0.276249	0.000000	0.000219
liveness	1610.0	0.491730	0.349100	0.021900	0.153000
loudness	1610.0	-6.971615	2.994003	-24.408000	-8.982500
speechiness	1610.0	0.069512	0.051631	0.023200	0.036500
tempo	1610.0	126.082033	29.233483	46.525000	107.390750
valence	1610.0	0.582165	0.231253	0.000000	0.404250
popularity	1610.0	20.788199	12.426859	0.000000	13.000000

	50%	75%	max
acousticness	0.18300	0.40375	0.994
danceability	0.45800	0.57800	0.887
energy	0.84850	0.94500	0.999
instrumentalness	0.01375	0.17900	0.996
liveness	0.37950	0.89375	0.998
loudness	-6.52300	-4.60875	-1.014
speechiness	0.05120	0.08660	0.624
tempo	124.40450	142.35575	216.304
valence	0.58300	0.77800	0.974
popularity	20.00000	27.00000	80.000

```
[8]: # Checking for the Null Values
df.isna().sum(axis=0)
```

```
[8]: name          0
      album        0
      release_date  0
      acousticness  0
      danceability  0
      energy        0
      instrumentalness  0
      liveness      0
      loudness      0
      speechiness   0
      tempo         0
      valence       0
      popularity    0
      dtype: int64
```

```
[9]: df.columns
```

```
[9]: Index(['name', 'album', 'release_date', 'acousticness', 'danceability',
          'energy', 'instrumentalness', 'liveness', 'loudness', 'speechiness',
          'tempo', 'valence', 'popularity'],
          dtype=object)
```

```
dtype='object')
```

```
[10]: a1 = []
a2 = list(np.sort(df['album'].unique()))

b1 = ["12 X 5", "December's Children", "England's Newest Hit Makers", "Get Yer_
↳Ya-Ya's Out", "Get Yer Ya-Ya's Out"]
b2 = ["12 x 5", "Decemberâ€s Children", "Englandâ€s Newest Hitmakers", "Get_
↳Yer Ya Yas Out", "Get Yer Ya-Ya's Out!"]

for i in np.sort(df['album'].unique()):
    a1.append(i.split(' (')[0].split(' - ')[0])

df['album'] = [i if i not in a2 else a1[a2.index(i)] for i in df['album']]
df['album'] = [i if i not in b2 else b1[b2.index(i)] for i in df['album']]
```

```
[11]: # Checking for the duplicatesNull Values
df[df[['name', 'album']].duplicated()]
```

```
[11]:
```

	name	album \
92	Start Me Up - Remastered 2021	Tattoo You
93	Hang Fire - Remastered 2021	Tattoo You
94	Slave - Remastered 2021	Tattoo You
95	Little T&A - Remastered 2021	Tattoo You
96	Black Limousine - Remastered 2021	Tattoo You
...	...	...
1592	I'm A King Bee	England's Newest Hit Makers
1593	Carol	England's Newest Hit Makers
1595	Can I Get A Witness	England's Newest Hit Makers
1596	You Can Make It If You Try	England's Newest Hit Makers
1597	Walking The Dog	England's Newest Hit Makers

	release_date	acousticness	danceability	energy	instrumentalness \
92	2021-10-22	0.0302	0.555	0.956	0.367000
93	2021-10-22	0.0136	0.421	0.927	0.079700
94	2021-10-22	0.1300	0.558	0.898	0.306000
95	2021-10-22	0.1560	0.497	0.953	0.042900
96	2021-10-22	0.1220	0.480	0.934	0.198000
...	...	...	...	...	...
1592	1964-05-30	0.0160	0.884	0.554	0.005620
1593	1964-05-30	0.1480	0.599	0.939	0.007470
1595	1964-05-30	0.3120	0.783	0.783	0.000000
1596	1964-05-30	0.2010	0.699	0.554	0.000051
1597	1964-05-30	0.3960	0.724	0.942	0.064800

	liveness	loudness	speechiness	tempo	valence	popularity
92	0.0753	-2.147	0.0577	121.752	0.933	12

93	0.2640	-1.734	0.0345	151.096	0.922	11
94	0.0562	-4.314	0.0515	133.338	0.853	11
95	0.6670	-1.978	0.0328	135.453	0.801	13
96	0.7010	-1.382	0.0715	108.894	0.825	9
...	...	...	...	...	...	...
1592	0.0786	-10.899	0.0636	105.536	0.884	0
1593	0.3200	-9.223	0.0390	88.645	0.967	1
1595	0.0763	-7.981	0.0741	97.018	0.842	0
1596	0.1070	-9.465	0.0529	102.508	0.582	0
1597	0.1160	-8.357	0.0354	125.331	0.967	0

[371 rows x 13 columns]

```
[12]: df.drop(df.index[df[df[['name', 'album']].duplicated()].index], inplace= True)
df[df[['name', 'album']].duplicated()]
```

```
[12]: Empty DataFrame
Columns: [name, album, release_date, acousticness, danceability, energy,
instrumentalness, liveness, loudness, speechiness, tempo, valence, popularity]
Index: []
```

```
[13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1239 entries, 0 to 1609
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                  1239 non-null   object
1   album                 1239 non-null   object
2   release_date          1239 non-null   datetime64[ns]
3   acousticness           1239 non-null   float64
4   danceability           1239 non-null   float64
5   energy                 1239 non-null   float64
6   instrumentalness       1239 non-null   float64
7   liveness               1239 non-null   float64
8   loudness               1239 non-null   float64
9   speechiness           1239 non-null   float64
10  tempo                  1239 non-null   float64
11  valence                1239 non-null   float64
12  popularity             1239 non-null   int64
dtypes: datetime64[ns](1), float64(9), int64(1), object(2)
memory usage: 135.5+ KB
```

```
[14]: df.head()
```

```
[14]:
```

		name	album	release_date	acousticness	\
0		Concert Intro Music - Live	Licked Live In NYC	2022-06-10	0.0824	
1		Street Fighting Man - Live	Licked Live In NYC	2022-06-10	0.4370	
2		Start Me Up - Live	Licked Live In NYC	2022-06-10	0.4160	
3		If You Can't Rock Me - Live	Licked Live In NYC	2022-06-10	0.5670	
4		Don't Stop - Live	Licked Live In NYC	2022-06-10	0.4000	

	danceability	energy	instrumentalness	liveness	loudness	speechiness	\
0	0.463	0.993	0.996000	0.932	-12.913	0.1100	
1	0.326	0.965	0.233000	0.961	-4.803	0.0759	
2	0.386	0.969	0.400000	0.956	-4.936	0.1150	
3	0.369	0.985	0.000107	0.895	-5.535	0.1930	
4	0.303	0.969	0.055900	0.966	-5.098	0.0930	

	tempo	valence	popularity
0	118.001	0.0302	33
1	131.455	0.3180	34
2	130.066	0.3130	34
3	132.994	0.1470	32
4	130.533	0.2060	32

```
feature_list = ['acousticness', 'danceability', 'energy', 'instrumentalness', 'liveness', 'loudness',
'speechiness', 'tempo', 'valence']
```

```
[15]: feature_list = ['acousticness', 'danceability', 'energy', 'instrumentalness',
↳ 'liveness', 'loudness', 'speechiness', 'tempo', 'valence']
features_df = pd.DataFrame(df[feature_list])
features_df.head()
```

```
[15]:
```

	acousticness	danceability	energy	instrumentalness	liveness	loudness	\
0	0.0824	0.463	0.993	0.996000	0.932	-12.913	
1	0.4370	0.326	0.965	0.233000	0.961	-4.803	
2	0.4160	0.386	0.969	0.400000	0.956	-4.936	
3	0.5670	0.369	0.985	0.000107	0.895	-5.535	
4	0.4000	0.303	0.969	0.055900	0.966	-5.098	

	speechiness	tempo	valence
0	0.1100	118.001	0.0302
1	0.0759	131.455	0.3180
2	0.1150	130.066	0.3130
3	0.1930	132.994	0.1470
4	0.0930	130.533	0.2060



## 1.2 Scaling And Decoposing

```
[16]: from sklearn.manifold import TSNE
      from sklearn.preprocessing import StandardScaler

      # scale the data for better results
      features_df = StandardScaler().fit_transform(features_df)

      features_df_tsne = TSNE(learning_rate=100).fit_transform(features_df)

      features_df_tsne
```

```
[16]: array([[ -19.618046, -33.72533 ],
             [-28.162659,  16.59559 ],
             [-30.906765, -22.013271],
             ...,
             [ 23.98327 , -10.512095],
             [ 39.72409 , -8.797641],
             [ 23.788965, -5.317569]], dtype=float32)
```

Decomposing the features as there are too many features for the clustering of the songs. Reducing the number of features to two where the features don;t really loose any of the reduced values

```
[17]: features_df = pd.DataFrame(data = features_df,
                                columns = feature_list)

      features_df.head()
```

```
[17]:
```

	acousticness	danceability	energy	instrumentalness	liveness	loudness	\
0	-0.711916	0.029653	1.054720	3.071082	1.162032	-2.045796	
1	0.853867	-0.939624	0.898093	0.263934	1.244453	0.672496	
2	0.761138	-0.515123	0.920468	0.878342	1.230242	0.627917	
3	1.427898	-0.635398	1.009969	-0.592902	1.056873	0.427146	
4	0.690488	-1.102349	0.920468	-0.387634	1.258664	0.573619	

	speechiness	tempo	valence
0	0.733260	-0.277206	-2.316806
1	0.075820	0.182247	-1.076932
2	0.829659	0.134813	-1.098472
3	2.333480	0.234803	-1.813619
4	0.405504	0.150761	-1.559440

## 2 DATA VISUALISATION

- Feature Distribution

```
[18]: fig = make_subplots(rows = 3,
                        cols = 3,
```

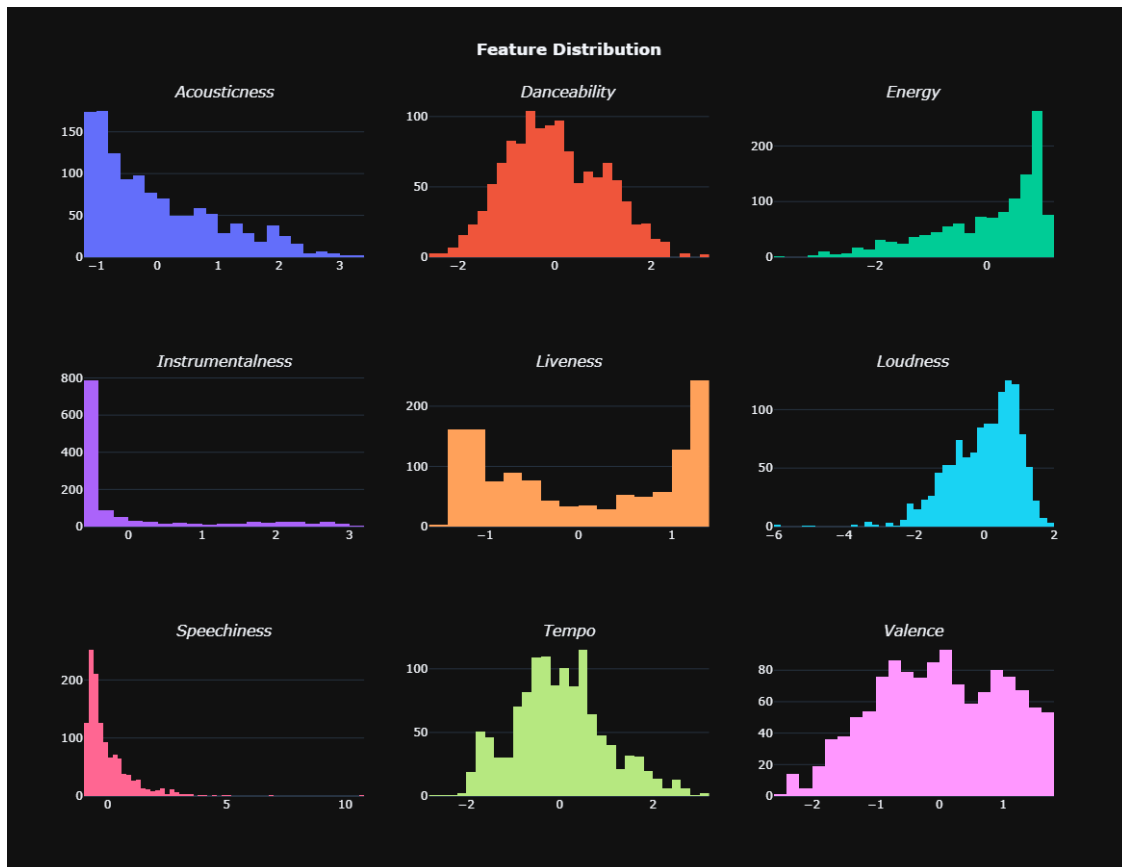
```

        subplot_titles=[f'<i>{i.title()}</i>' for i in feature_list])
l = len(feature_list)
for i in range(l):
    fig.add_trace(go.Histogram(x = features_df[feature_list[i]]),
                    col = i % 3 + 1,
                    row = i // 3 + 1)

fig.update_layout(height = 900,
                  width = 900,
                  title_text = '<b>Feature Distribution</b>',
                  template = 'plotly_dark',
                  title_x = 0.5,
                  showlegend=False)


fig.show()

```



- Violin Plot for Different Attributes

```
[19]: px.violin(data_frame= features_df,
               y = feature_list,
               color_discrete_sequence = ['turquoise'],
               template = 'plotly_dark',
               title = '<b>Violin Plot for Different Attributes')
```



```
[20]: def outliers(feature_list, dataset) :

    l = len(feature_list)
    row = 3
    col = l//row
    fig = make_subplots(rows = row,
                        cols = col,
                        subplot_titles=[f'<i>{i.title()}' for i in_
↪feature_list])
    for i in range(l):
        fig.add_trace(go.Box(x = dataset[feature_list[i]]),
                        col = i % col + 1,
                        row = i // row +1)

    fig.update_layout(height = 900,
                      width = 900,
                      title_text = '<b>Feature Distribution')

    fig.update_layout(template = 'plotly_dark',
                      title_x = 0.5)
    fig.update_layout(showlegend=False)

    fig.show()

outliers(feature_list, df )
```



We won't be removing outliers as outlying part of one feature is major part of another as can be seen from Instrumentalness and Speechiness


- Heat Map for Different Attributes

```
[21]: fig = px.imshow(features_df.corr(),
    text_auto = True,
    height = 1000,
    template = 'plotly_dark',
    title = '<b>Heat Map for Different Attributes')
fig.show()
```



- No Strong correlation found except for between Loudness and Energy
- Energy Versus Loudness


```
[22]: fig = px.scatter(features_df,
    x = 'loudness',
    y = 'energy',
    color = 'loudness',
    color_continuous_scale = 'tealrose',
    template = 'plotly_dark',
    title = f'<b>Energy Versus Loudness')
fig.show()
```



- Songs Released by Year

```
[23]: df['year'] = pd.DatetimeIndex(df['release_date']).year

fig = px.bar(data_frame = df.groupby('year',as_index=False).count().
    ↪sort_values(by='name',ascending=False).sort_values(by='year'),
            x = 'year',
            y = 'name',
            labels = {'name':'Total songs'},
            color_discrete_sequence = ['silver'],
            template = 'plotly_dark',
            title = '<b>Songs Released by Year')
fig.update_layout(hovermode = 'x',
                  title_x = 0.5)
```




- Top 25 songs in Spotify

```
[24]: fig = px.line(df.sort_values(by='popularity',
    ascending=False).head(25),
            x = 'name',
            y = 'popularity',
            hover_data = ['album'],
```

```

color_discrete_sequence = ['silver'],
template = 'plotly_dark',
markers = True,
title = '<b> Top 25 songs in Spotify',
)
fig.show()

```




- Recommended Album based on Average of Popularity of Songs

```


[25]: by_mean = df.groupby('album')['popularity'].mean().sort_values(ascending =
      ↪False).head(15)
fig = px.bar(by_mean,
             y = 'popularity',
             color = 'popularity',
             color_continuous_scale = 'plasma',
             template = 'plotly_dark',
             title = f'<b>Recommended Album based on Average of Popularity of Songs')
fig.show()

```



- Recommended Album based on Most Popular Song


```
[26]: by_max = df.groupby('album')['popularity'].max().sort_values(ascending = False).
      ↪head(15)
fig = px.bar(by_max,
             y = 'popularity',
             color = 'popularity',
             color_continuous_scale = 'plasma',
             template = 'plotly_dark',
             title = f'<b>Recommended Album based on Most Popular Song')
fig.show()
```



- Song Count

```
[27]: by_count = df.groupby('album')['popularity'].count().sort_values(ascending = 
      ↪False)
fig = px.bar(by_count,
             y = 'popularity',
             color = 'popularity',
             color_continuous_scale = 'plasma',
             template = 'plotly_dark',
             height = 800,
             title = f'<b>Count of Songs in Album')
fig.show()
```





Plotting various features against the Popularity of the songs

```
[28]: # Initialising the plot variable as str
fig1 = ''
fig2 = ''
fig3 = ''
fig4 = ''
fig5 = ''
fig6 = ''
fig7 = ''
fig8 = ''
fig9 = ''

# Storing the variables in list figures
figures = [fig1, fig2, fig3, fig4, fig5, fig6, fig7, fig8, fig9]

# Creating plots using for loop
for i in range(9):
    figures[i] = px.scatter(features_df,
                           x = feature_list[i],
                           y = df['popularity'],
                           color = feature_list[i],
                           color_continuous_scale = 'plasma',
```

```

trendline='ols')

# Storing the headers for the Names of Plots in a list
figure_names = []
for i in feature_list:
    figure_names.append(f'<i>{i.title()} VS Popularity')

# Initialising the trace list and then storing them into another list for
↳ looping
figure1_traces = []
figure2_traces = []
figure3_traces = []
figure4_traces = []
figure5_traces = []
figure6_traces = []
figure7_traces = []
figure8_traces = []
figure9_traces = []

figure_traces = [figure1_traces , figure2_traces, figure3_traces,
↳ figure4_traces, figure5_traces, figure6_traces, figure7_traces,
↳ figure8_traces, figure9_traces]

# Taking the elements from the plots and using them to create traces for the
↳ subplots
for i in range(9):
    for trace in range(len(figures[i]["data"])):
        figures[i]["data"][trace]['showlegend'] = False
        figure_traces[i].append(figures[i]["data"][trace])

# Creating the subplot Figure
comparison_figure = make_subplots(rows = 9, cols = 1, subplot_titles =
↳ figure_names)

# Adding the traces to the figure
for i in range(9):
    for traces in figure_traces[i]:
        comparison_figure.append_trace(traces, row = i+1, col = 1)

comparison_figure.update_layout(height = 2500, width = 1000, title_text =
↳ "Comaprison Graphs", title_font_size = 25, template = 'plotly_dark')

comparison_figure.show()

```

Comaprison Graphs



Through the trendline in the graphs, it can be seen that Acoustiness, Danceability, Valence, Loudness is liked more at the higher levels while the rest are less popular the higher their value or no effect at all.

## 2.1 Training Models for predicting Popularity and testing

```
[29]: from sklearn.linear_model import LinearRegression
      from sklearn.tree import DecisionTreeRegressor
      from sklearn.ensemble import RandomForestRegressor
      from xgboost import XGBRegressor
      from sklearn.model_selection import train_test_split
```

```
[30]: # separate the data to training and testing
      X= features_df # all the features except DV
      y = df["popularity"] # the DV
      X_train, X_test, y_train,y_test=train_test_split(X,y,
                                                         test_size=0.2,
                                                         random_state=0)

      # save as np.array
      X_train = np.array(X_train)
      X_test = np.array(X_test)
      y_train = np.array(y_train)
      y_test = np.array(y_test)
```

```
[31]: # create a linear regression, random forest & decision tree object
      model_regression = LinearRegression()
      model_random_forest = RandomForestRegressor()
      model_decision_tree = DecisionTreeRegressor()
      model_xgboost = XGBRegressor()
```

```
[32]: model_regression.fit(X_train, y_train)
```

```
[32]: LinearRegression()
```

```
[33]: model_random_forest.fit(X_train, y_train)
```

```
[33]: RandomForestRegressor()
```

```
[34]: model_decision_tree.fit(X_train, y_train)
```

```
[34]: DecisionTreeRegressor()
```

```
[35]: model_xgboost.fit(X_train, y_train)
```

```
[35]: XGBRegressor(base_score=0.5, booster='gbtree', callbacks=None,
                 colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
                 early_stopping_rounds=None, enable_categorical=False,
                 eval_metric=None, gamma=0, gpu_id=-1, grow_policy='depthwise',
                 importance_type=None, interaction_constraints='',
                 learning_rate=0.300000012, max_bin=256, max_cat_to_onehot=4,
                 max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
                 missing=nan, monotone_constraints='()', n_estimators=100, n_jobs=0,
                 num_parallel_tree=1, predictor='auto', random_state=0, reg_alpha=0,
                 reg_lambda=1, ...)
```

```
[36]: # estimate the  $R^2$  score on training and testing data
      # (1) Linear regression
      print(model_regression.score(X_train,y_train))
      print(model_regression.score(X_test,y_test))
```

```
0.19396651223160155
0.16926201652752848
```

```
[37]: # (3) Decision Tree
      print(model_decision_tree.score(X_train,y_train))
      print(model_decision_tree.score(X_test,y_test))
```

```
1.0
-0.725561921406735
```

```
[38]: # (2) Random Forest
      print(model_random_forest.score(X_train,y_train))
      print(model_random_forest.score(X_test,y_test))
```

```
0.8757247990646652
0.08522159574174182
```

```
[39]: # (4) XGBOOST
      print(model_xgboost.score(X_train,y_train))
      print(model_xgboost.score(X_test,y_test))
```

```
0.9967689107870302
-0.08023682210155036
```

All the models have very low accuracy score

## 2.2 Clustering

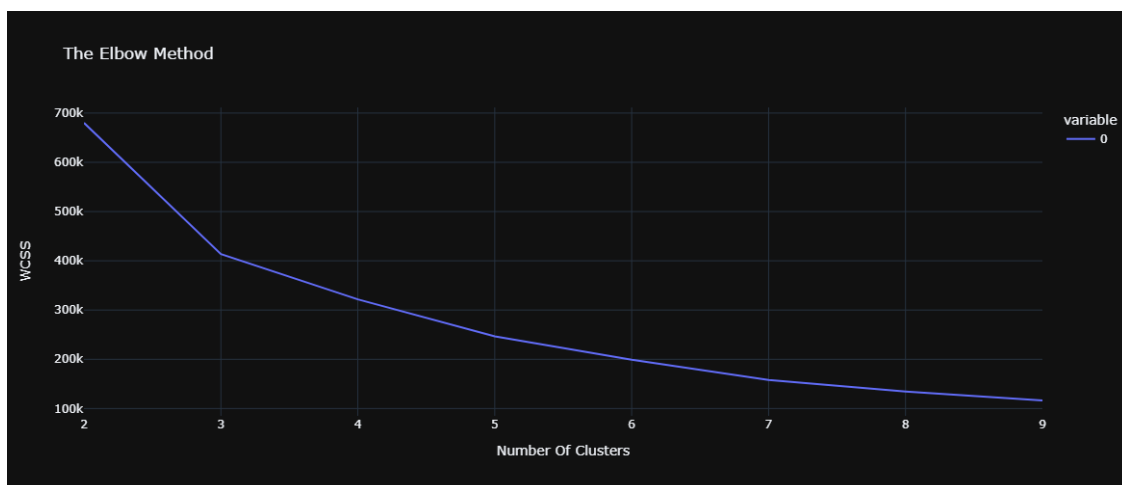
```
[40]: color = ['red', 'green', 'yellow', 'blue', 'magenta', 'cyan', 'white',
               ↪ 'gray', 'bright_red', 'bright_green', 'bright_yellow', 'bright_blue',
               ↪ 'bright_magenta', 'bright_cyan', 'bright_white']
```

### 2.2.1 K-Means Clustering

```
[41]: feature_arr = features_df_tsne
```

```
[42]: # ----- k-means Clustering -----  
  
# Determining the cluster size  
score_list = []  
for i in range(2,10):  
    kmeans_model = KMeans(n_clusters=i,  
                           init = 'k-means++',  
                           random_state=42).fit(feature_arr)  
  
    score_list.append(kmeans_model.inertia_)  
preds = kmeans_model.predict(feature_arr)
```

```
[43]: # Visualization of different cluster size performances  
  
fig = px.line(pd.DataFrame(score_list,  
                           index = range(2, 10)),  
              title = 'The Elbow Method')  
fig.update_layout(xaxis_title = 'Number Of Clusters',  
                  yaxis_title = 'WCSS',  
                  template = 'plotly_dark',  
                  height = 500,  
                  width = 500)  
fig.show()
```



```
[44]: cluster_i = 8
```

```
[45]: # Training and Predicting
k_pred = KMeans(n_clusters=cluster_i,
                init='k-means++',
                max_iter=300,
                random_state=42).fit_predict(feature_arr)
k_pred
```

```
[45]: array([3, 4, 3, ..., 1, 1, 1], dtype=int32)
```

```
[46]: # Initialising the plot variable as str
k_fig1 = ""
k_fig2 = ""
k_fig3 = ""
k_fig4 = ""
k_fig5 = ""
k_fig6 = ""
k_fig7 = ""
k_fig8 = ""

# Storing the variables in list figures
k_figures = [k_fig1, k_fig2, k_fig3, k_fig4, k_fig5, k_fig6, k_fig7, k_fig8]

# Creating cluster graphs using for loop
for i in range(8):
    k_figures[i] = px.scatter(x = feature_arr[k_pred == i, 0],
                             y = feature_arr[k_pred == i, 1],
                             title = f'Cluster{i}')
    k_figures[i].update_traces(marker=dict(color=color[i]))

# Initialising the trace list and then storing them into another list for
↳ looping
k_figure1_traces = []
k_figure2_traces = []
k_figure3_traces = []
k_figure4_traces = []
k_figure5_traces = []
k_figure6_traces = []
k_figure7_traces = []
k_figure8_traces = []

k_figures_traces = [k_figure1_traces, k_figure2_traces, k_figure3_traces,
↳ k_figure4_traces, k_figure5_traces, k_figure6_traces, k_figure7_traces,
↳ k_figure8_traces]
```

```

# Taking the elements from the plots and using them to create traces for the
↳subplots
for i in range(8):
    for trace in range(len(k_figures[i]["data"])):
        k_figures[i]["data"][trace]['showlegend'] = False
        k_figures_traces[i].append(k_figures[i]["data"][trace])

# Creting the subplot Figure
k_means_fig = make_subplots(rows = 1, cols = 1)

# Adding the traces to the figure
for i in range(8):
    for traces in k_figures_traces[i]:
        k_means_fig.append_trace(traces, row = 1, col = 1)

k_means_fig.update_layout(height = 800, width = 1000, title_text = "K-Mean Song
↳Clusters", title_font_size = 25, template = 'plotly_dark')

k_means_fig.show()

```






### 3 Heirchical Clustering

```
[47]: data = df['popularity'].to_numpy()
```

```
[48]: dendrogram=sch.dendrogram(sch.linkage(feature_arr,  
                                           method='ward'))
```



```
[49]: hrc_clusters = 8
```

```
[50]: hrc=AgglomerativeClustering(n_clusters=hrc_clusters,  
                                  metric='euclidean',  
                                  linkage='ward')
```

```
[51]: pred_hrc=hrc.fit_predict(feature_arr)  
pred_hrc
```

```
[51]: array([6, 0, 6, ..., 3, 3, 3])
```

```
[52]: # Initialising the plot variable as str  
hrc_fig1 = ""  
hrc_fig2 = ""  
hrc_fig3 = ""  
hrc_fig4 = ""  
hrc_fig5 = ""  
hrc_fig6 = ""
```

```

hrc_fig7 = ""
hrc_fig8 = ""
hrc_fig9 = ""
hrc_fig10 = ""

# Storing the variables in list figures
hrc_figures = [hrc_fig1, hrc_fig2, hrc_fig3, hrc_fig4, hrc_fig5, hrc_fig6,
↳hrc_fig7, hrc_fig8, hrc_fig9, hrc_fig10]

# Creating cluster graphs using for loop
for i in range(hrc_clusters):
    hrc_figures[i] = px.scatter(x = feature_arr[pred_hrc == i, 0],
                               y = feature_arr[pred_hrc == i, 1],
                               title = f'Cluster{i}')
    hrc_figures[i].update_traces(marker=dict(color=color[i]))

# Initialising the trace list and then storing them into another list for
↳looping
hrc_figure1_traces = []
hrc_figure2_traces = []
hrc_figure3_traces = []
hrc_figure4_traces = []
hrc_figure5_traces = []
hrc_figure6_traces = []
hrc_figure7_traces = []
hrc_figure8_traces = []
hrc_figure9_traces = []
hrc_figure10_traces = []

hrc_figures_traces = [hrc_figure1_traces, hrc_figure2_traces,
↳hrc_figure3_traces, hrc_figure4_traces, hrc_figure5_traces,
↳hrc_figure6_traces, hrc_figure7_traces, hrc_figure8_traces,
↳hrc_figure9_traces, hrc_figure10_traces]

# Taking the elements from the plots and using them to create traces for the
↳subplots
for i in range(hrc_clusters):
    for trace in range(len(hrc_figures[i]["data"])):
        hrc_figures[i]["data"][trace]['showlegend'] = False
        hrc_figures_traces[i].append(hrc_figures[i]["data"][trace])

# Creting the subplot Figure
hrc_cluster_figure = make_subplots(rows = 1, cols = 1)

```


```

# Adding the traces to the figure
for i in range(hrc_clusters):
    for traces in hrc_figures_traces[i]:
        hrc_cluster_figure.append_trace(traces, row = 1, col = 1)

hrc_cluster_figure.update_layout(height = 800, width = 1000, title_text = "HRC_
↳Song Clusters", title_font_size = 25, template = 'plotly_dark')

hrc_cluster_figure.show()

```



[ ]:

4 Thank You